*one team. one route.*

*five developers. five routes.*

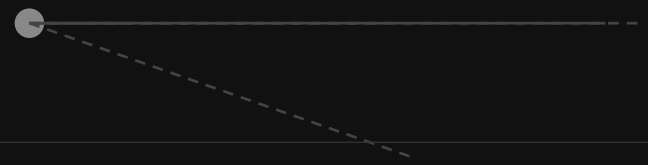

# AI is GPS. We're driving.

## Most of us have the volume off.

*Don't let AI work for you. Work better with AI.*

# Same tools. Same team. Completely different results.

**what we expect**

We use AI → we save time

We review code → consistent output

Same tools → same results

**what actually happens**

Everyone has their own "AI style"

No shared workflow → no predictability

Copy → paste → ship → debug → repeat

No measurement → no improvement

***“We don’t have an AI problem. We have a process problem.”***

# One route. Everyone follows it. Everyone arrives.

### Shared workflow

One process. Every task. Every person.

### Shared context

CLAUDE.md + .cursorrules per project.

### Clear tool roles

Claude plans · Cursor builds · Copilot assists.

### Feedback loop

Measure what ships. Learn. Adjust.

PR cycles ↓ Bug rate ↓ Onboarding time ↓ **Predictability ↑**

*“Consistency beats individual brilliance.”*

# Right instrument. Right moment.

| Tool | Role | Best for |
|---|---|---|
| **Claude** | Long-range planner | Architecture · Planning · Complex reasoning · Long documents |
| **Cursor** | Turn-by-turn navigation | In-file editing · Multi-file refactoring · Codebase-aware changes |
| **Copilot** | Heads-up display | Inline autocomplete · PR summaries · Fast completions |
| **ChatGPT** | Offline map | Quick Q&A · Exploration · No project context |

← Inline / Fast / Low context

Deep / Planned / High context →

Copilot · ChatGPT · Cursor · Claude

*“One instrument per task. Don’t drive staring at three screens simultaneously.”*

# GPS told someone to drive into a lake. They followed it.

**BLIND**

open AI tool → generate → commit

compiles? ship it.

PR fails → back to start

× PR rejected / prod bug

*3 days*

**SMART**

plan with Claude first

generate → read every line

"what did AI miss?" → test it

commit what you can explain

✓ **Merged**

*1 day*

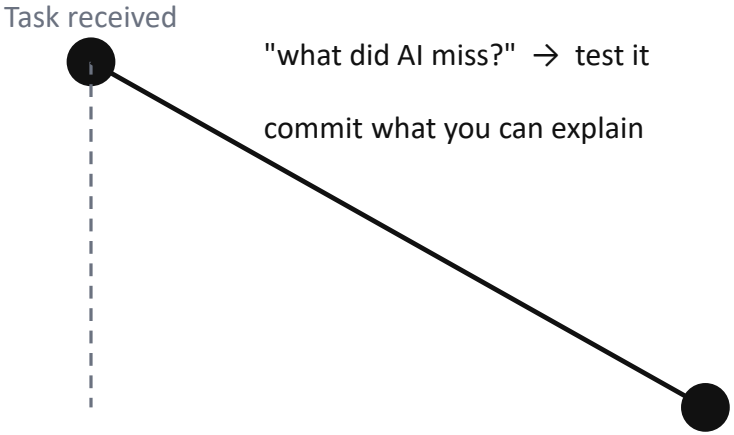


*"Same AI. Same codebase. One day versus three."*

# You don’t drive without setting the destination first.

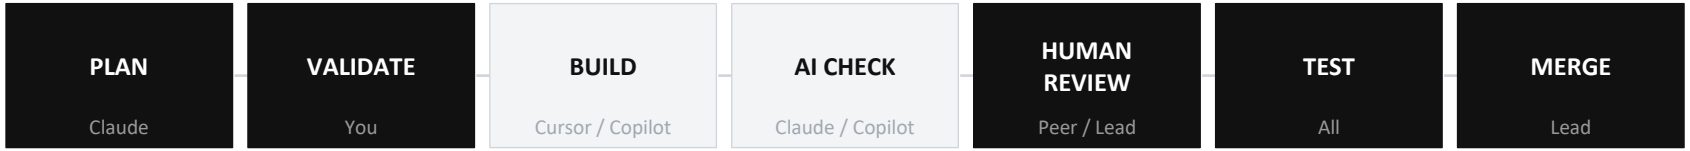


■ Non-negotiable □ Recommended

| | | |
|---|---|---|
| **PLAN** | Claude | Break down task, identify edge cases, confirm approach |
| **VALIDATE** | You | Confirm plan is correct before a line is written |
| **BUILD** | Cursor / Copilot | Implementation with context from the plan |
| **AI CHECK** | Claude | Logic gaps, edge cases, pattern check |
| **HUMAN REVIEW** | Peer / Lead | Business logic, architecture, security |
| **TEST** | All | Unit + edge cases alongside the code, not after |
| **MERGE** | Lead / Peer | Explicit approval — lead for core, peer for rest |

# Recalculating. Five ways to stay on course.

| Control | When to use | What it does |
|---|---|---|
| **Continue** | Claude paused mid-task | Picks up exactly where it left off — no drift. Re-prompting restarts reasoning and can drift. |
| **Rewind** | Wrong direction taken | Cuts the wrong path cleanly. Re-prompt with the correct target. |
| **/clear** | Switching tasks entirely | Removes stale context. Fresh map for a new destination. |
| **/compact** | Context window filling up | Claude summarises progress. Recovers ~60% context. Compact before it fills — not after. |
| **Subagent** | Repeatable specialist task | /review triggers code-reviewer.md — same team standards, every PR, automatically. |

*"Re-prompting is not the same as recalculating."*

# Wrong address in. Wrong place out. Every time.

**Context + Task + Constraints + Format = Useful output**

WEAK PROMPT

```
Fix this function
```

Here's a possible fix... You might also consider... Depending on your framework... There are several approaches...

STRONG PROMPT

```
Node.js + Express + PostgreSQL.
/api/users/:id → 500 when user missing.
Fix: return 404 { error: "user_not_found" }
No JWT/signature changes. Code only.
```

```
try {
  const user = await db.findById(id);
  if (!user) return res.status(404)
    .json({ error: "user_not_found" });
}
```

1. Name the file / function — don't describe it
2. State what you've already tried
3. Specify format: code only · bullet list · max 3 options
4. One task per prompt — chained prompts drift
5. Add constraints: what NOT to change

# Don’t let the signal drop mid-journey.

**HIGH WASTE**

— Pasting 500 lines for a 3-line question

— Re-explaining context every session

— Vague questions → 3 clarifying rounds

Context: 82% used

**STRONG SIGNAL**

+ /compact before context fills → 60% back

+ /clear when switching tasks

+ Reference line numbers, not whole files

Context: 28% used

| Model | Use for | Cost |
|---|---|---|
| **Haiku** | Repetitive tasks · test gen · summaries | Low |
| **Sonnet** | Most dev tasks — best quality / cost balance | Medium |
| **Opus** | Architecture · complex reasoning only | High |

# AI knows the route before you say a word.

```
project/
  CLAUDE.md
  .cursorrules
  .claude/
    agents/
      code-reviewer.md
      test-generator.md
    commands/
      review.md
```

**CLAUDE.md**

Claude's route preference
Stack · patterns · off-limits zones

**.cursorrules**

Cursor's driving style
Coding conventions + what to avoid

**agents/**

Repeatable specialist tasks
/review · /gen-tests · /docs

**commands/**

One-command team standards
Same quality, every time

Minimum CLAUDE.md — copy this for every project:

```
Stack: Node.js 20, React 18, PostgreSQL 14, AWS
Style: ESLint Airbnb, async/await, no raw SQL (use Knex)
```

# Know exactly what your AI can see — and what it’s guessing.

## Claude

Sees: What you paste / attach / reference

*Blind: Everything else in your codebase*

## Cursor

Sees: Open + indexed files in the project

*Blind: Runtime state · excluded files · secrets*

## Copilot

Sees: Current file + adjacent open files

*Blind: Your whole project context + standards*

## ChatGPT

Sees: This conversation only

*Blind: Your entire codebase + stack*

**Signs AI is guessing:**

— Imports a function that doesn't exist

— Suggests a package you already have

— Recommends a pattern you don't use

# GPS says turn right. You still look at the road.

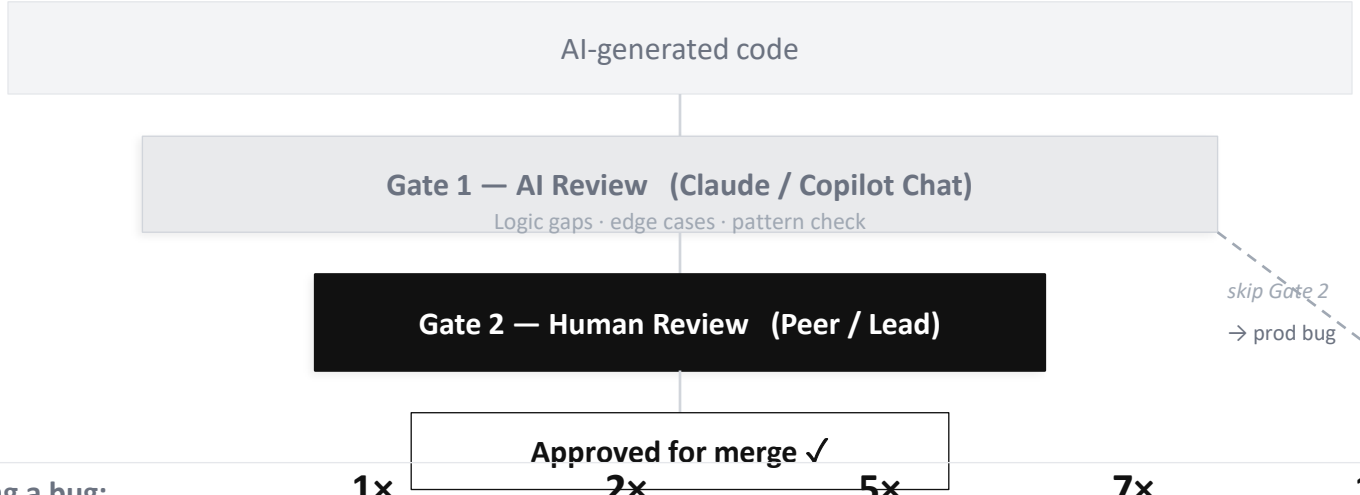


Cost of finding a bug:

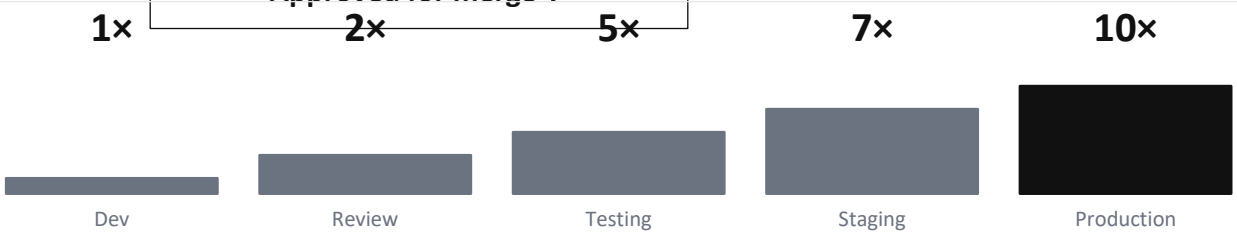

# You still drive. GPS is not the driver.

| | |
|---|---|
| **GREEN**<br>Full delegate | Boilerplate · CRUD scaffolding · test stubs · PR descriptions · docs |
| **YELLOW**<br>Collaborate | Business logic · integrations · refactoring · query optimisation |
| **AMBER**<br>Human-led | System design · API contracts · data models · critical flows |
| **RED**<br>Never delegate | Auth · encryption · PII · production incidents · architecture decisions |

**Ask before every task:**

Can the output be verified quickly? → higher delegation

Does it need context AI doesn't have? → lower delegation

Could wrong output reach production silently? → never delegate

# GPS or not — road rules still apply.

**Non-negotiables**

1. All AI-generated code gets human review — no exceptions
2. No credentials, API keys, or client data in any AI prompt
3. Auth, encryption, PII handling: manual development only
4. "Can't explain it → don't commit it" — escalate
5. No AI agent with write access to production systems

**DO ✓**

Review and understand every AI output before committing

Use CLAUDE.md and .cursorrules on every active project

Tag PRs: ai-assisted + tool used

Log AI wins and failures in sprint retros

**DON'T ✗**

Deploy without validation — even if it compiles

Skip human review because "AI already checked it"

Let AI make architecture decisions without team sign-off

Paste client data or business logic into public AI tools

***"AI suggests something → you can't explain it → escalate. Don't commit."***

# You can’t improve a route you don’t measure.

| PR cycle time | Rework rate | AI-assist rate | Defect escape |
|---|---|---|---|
| Git: open → merge (hours) | PRs needing >1 review (%) | Self-reported per sprint | Bugs found post-merge |

| Who | Action | When |
|---|---|---|
| **All devs** | Create CLAUDE.md for top 2 active projects | This sprint |
| **All devs** | Log AI-assisted tasks in sprint notes | This sprint |
| **Tech leads** | Add ai-assisted tag to PR template | This week |
| **Architects** | Define AI off-limits zones per project | This week |

1. AI increases speed — not correctness
2. Discipline makes the difference — not the tool
3. Consistency beats individual brilliance
4. Measure it. Improve it. Don't assume it.

*“The best AI workflow is the one your whole team actually follows.”*